

Mojca Kompara, Miha Meglič, Andraž Sovinec

The abstract goes here. The abstract goes here. The abstract goes here. The abstract goes here. The abstract
goes here. The abstract goes here. The abstract goes here. The abstract goes here. The abstract goes here.
The abstract goes here. The abstract goes here. The abstract goes here. The abstract goes here. The abstract
goes here. The abstract goes here. The abstract goes here. The abstract goes here. The abstract goes here.
The abstract goes here. The abstract goes here. The abstract goes here. The abstract goes here. The abstract
goes here. The abstract goes here. The abstract goes here. The abstract goes here.

Page | 8

LLM, chatbot, RAG

Introduction

The primary motivation for this project is to reduce the burden on student office staff and provide students with instant, reliable answers at any time. the RAG approach allows us to work with a relatively small corpus of institutional documents while still having the generative capabilities of modern large language models.

UL FS has already implemented a chatbot for student inquiries, but it relies on a simple and short inputs and cannot handle complex institutional queries (<https://www.fs.uni-lj.si/en/>).

To address complex queries over large document collections, Lee [2] proposes a RAG framework that decomposes multihop questions into single-hop subquestions using an LLM. Closer to home, the UL FS example chatbot demonstrates practical institutional use for student queries, inspiring our FRI-specific adaptation. These works validate RAG’s effectiveness for administrative chatbots (<https://www.fs.uni-lj.si/en/>).

Proposed Data

For this project, we will use web-accessible data, for example UL FRI web pages (<https://www.fri.uni-lj.si/>), timetables (<https://urnik.fri.uni-lj.si/timetable/>), and documents.

References

- [1] Deep Shah. College rag chatbot. https://github.com/DeepShah1406/College_RAG_Chatbot, 2024. Accessed: March 20, 2026.
- [2] Seokgi Lee. Transforming questions and documents for semantically aligned retrieval-augmented generation. *arXiv preprint arXiv:2508.09755*, 2025.